

```
#importing the libraries
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
sns.set()
import warnings
warnings.filterwarnings('ignore')

#loading the dataset
df = pd.read_csv('home_data.csv')
df.head()
```

	id	date	price	bedrooms	bathrooms
sqft_living \					
0	7129300520	20141013T000000	221900	3	1.00
1180					
1	6414100192	20141209T000000	538000	3	2.25
2570					
2	5631500400	20150225T000000	180000	2	1.00
770					
3	2487200875	20141209T000000	604000	4	3.00
1960					
4	1954400510	20150218T000000	510000	3	2.00
1680					

	sqft_lot	floors	waterfront	view	...	grade	sqft_above
sqft_basement \							
0	5650	1.0	0	0	...	7	1180
0							
1	7242	2.0	0	0	...	7	2170
400							
2	10000	1.0	0	0	...	6	770
0							
3	5000	1.0	0	0	...	7	1050
910							
4	8080	1.0	0	0	...	8	1680
0							

	yr_built	yr_renovated	zipcode	lat	long	sqft_living15	\
0	1955	0	98178	47.5112	-122.257	1340	
1	1951	1991	98125	47.7210	-122.319	1690	
2	1933	0	98028	47.7379	-122.233	2720	
3	1965	0	98136	47.5208	-122.393	1360	
4	1987	0	98074	47.6168	-122.045	1800	

	sqft_lot15
0	5650
1	7639
2	8062

```
3      5000
4      7503
```

```
[5 rows x 21 columns]
```

Some Numerical Information about the Data

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21613 entries, 0 to 21612
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    21613 non-null  int64
1   date                 21613 non-null  object
2   price                21613 non-null  int64
3   bedrooms             21613 non-null  int64
4   bathrooms            21613 non-null  float64
5   sqft_living          21613 non-null  int64
6   sqft_lot             21613 non-null  int64
7   floors               21613 non-null  float64
8   waterfront           21613 non-null  int64
9   view                 21613 non-null  int64
10  condition             21613 non-null  int64
11  grade                21613 non-null  int64
12  sqft_above           21613 non-null  int64
13  sqft_basement        21613 non-null  int64
14  yr_built              21613 non-null  int64
15  yr_renovated          21613 non-null  int64
16  zipcode              21613 non-null  int64
17  lat                  21613 non-null  float64
18  long                 21613 non-null  float64
19  sqft_living15         21613 non-null  int64
20  sqft_lot15           21613 non-null  int64
dtypes: float64(4), int64(16), object(1)
memory usage: 3.5+ MB
```

```
df.nunique()
```

```
id          21436
date         372
price       4032
bedrooms     13
bathrooms    30
sqft_living 1038
sqft_lot     9782
floors        6
waterfront   2
view         5
```

condition	5
grade	12
sqft_above	946
sqft_basement	306
yr_built	116
yr_renovated	70
zipcode	70
lat	5034
long	752
sqft_living15	777
sqft_lot15	8689
dtype:	int64

Data Cleaning

```
# Redefine date column to year and month
df['date'] = df['date'].str[2:6]

# Define a function to change type of bathroom column to int
def round_x(x):
    return (int(x+0.25))

df['bathrooms'] = df['bathrooms'].apply(lambda x : round_x(x))

# Define a function for reduce uniques of yr_renovated column to two value (0,1)
def year(x):
    if x == 0 :
        return 0
    else :
        return 1

df['yr_renovated'] = df['yr_renovated'].apply(lambda x : year(x))

# Reduce unique values of floor column
df['floors'] = df['floors'].apply(lambda x : 3 if x == 3.5 else x)

# Drop outliers
df = df[df['price'] < 2000000]
```

Data Visualization

```
# Define list of Continuous columns Names
continuous = ['price', 'sqft_living']

# Define a function to Capitalize the first element of string and remove '_' character
def title(name):
    return (' '.join(word.capitalize() for word in name.split('_')))
```

```

# Distribution of Categorical Features
def plot_continuous_distribution(df, column):

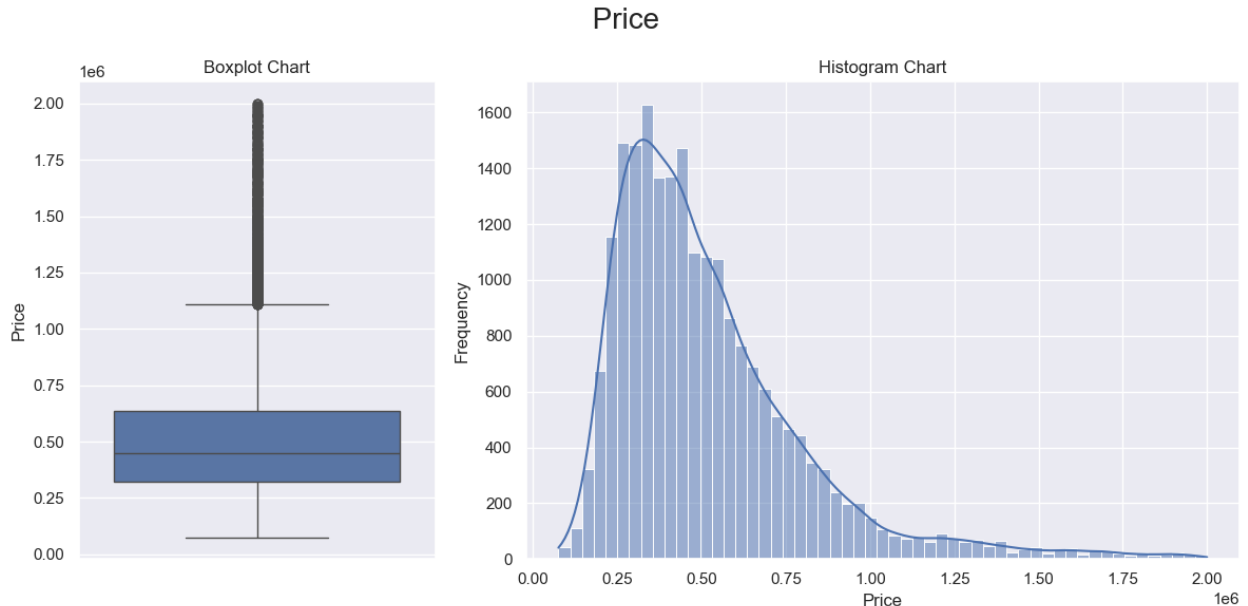
    width_ratios = [2, 4]
    gridspec_kw = {'width_ratios':width_ratios}
    fig, ax = plt.subplots(1, 2, figsize=(12, 6), gridspec_kw =
gridspec_kw)
    fig.suptitle(f' {title(column)} ', fontsize=20)


    sns.boxplot(df[column], ax=ax[0])
    ax[0].set_title('Boxplot Chart')
    ax[0].set_ylabel(title(column))

    sns.histplot(x = df[column], kde=True, ax=ax[1], multiple =
'stack', bins=55)
    ax[1].set_title('Histogram Chart')
    ax[1].set_ylabel('Frequency')
    ax[1].set_xlabel(title(column))

    plt.tight_layout()
    plt.show()
for conti in continuous :
    plot_continuous_distribution(df, conti)

```





```
# Define list of Categorical columns Names
categorical = ['waterfront', 'yr_renovated', 'floors', 'condition',
               'view']

# distribution of categorical features

def plot_categorical_distribution(df, column):
    fig, ax = plt.subplots(1, 2, figsize=(10, 6))
    fig.suptitle(f' {title(column)} ', fontsize=20)

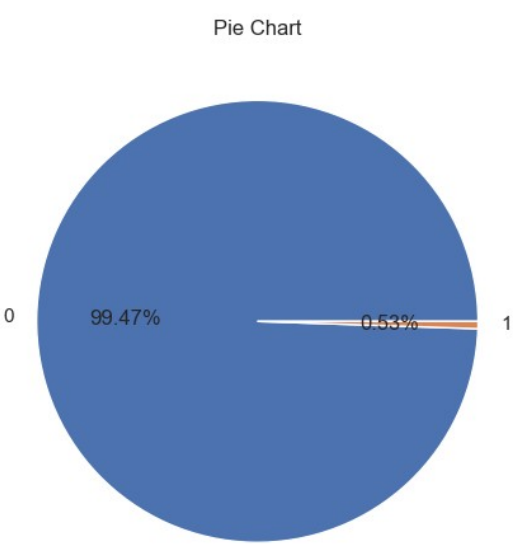

    sns.barplot(df[column].value_counts(), ax=ax[0], palette='deep')
    ax[0].set_title('Bar Chart')
    ax[0].set_xlabel(title(column))
    ax[0].set_ylabel('Frequency')

    df[column].value_counts().plot(kind='pie', autopct="%.2f%%",
    ax=ax[1])
    ax[1].set_title('Pie Chart')
    ax[1].set_ylabel(None)

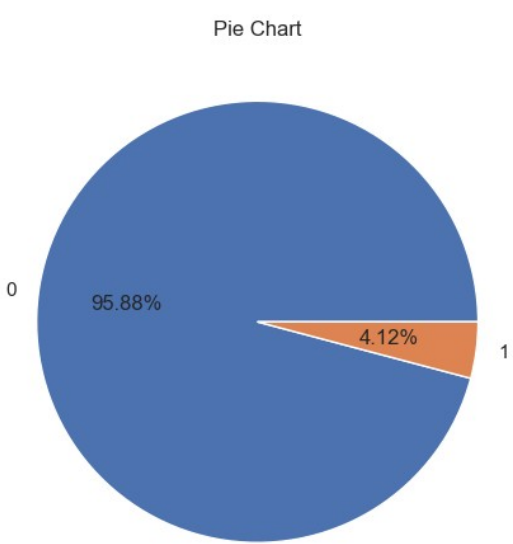
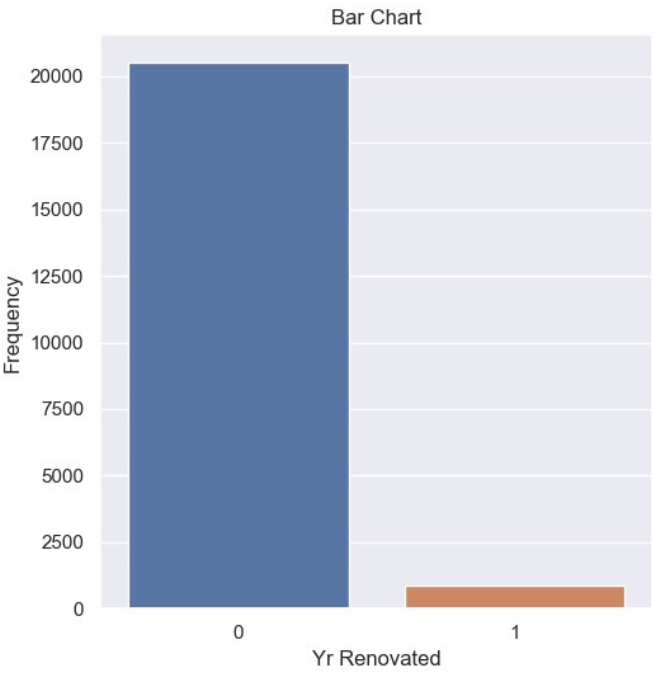
    plt.tight_layout()
    plt.show()

for cat in categorical:
    plot_categorical_distribution(df, cat)
```



Waterfront




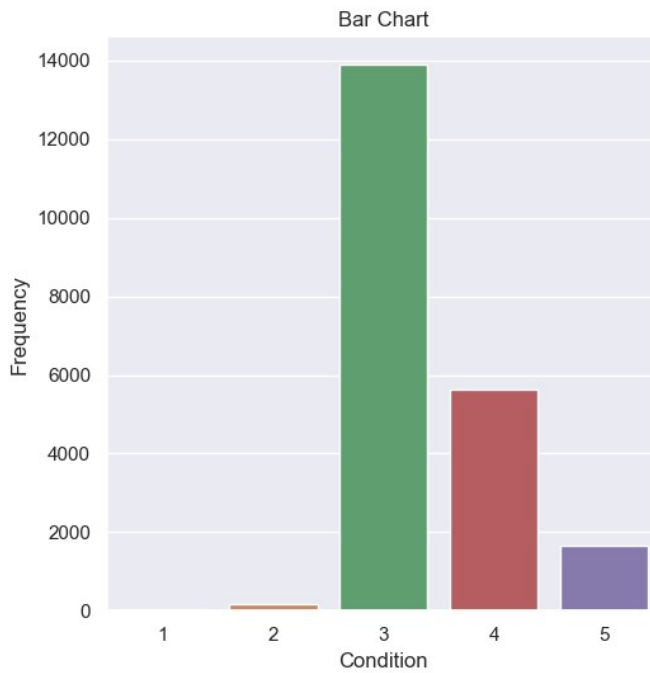
Yr Renovated



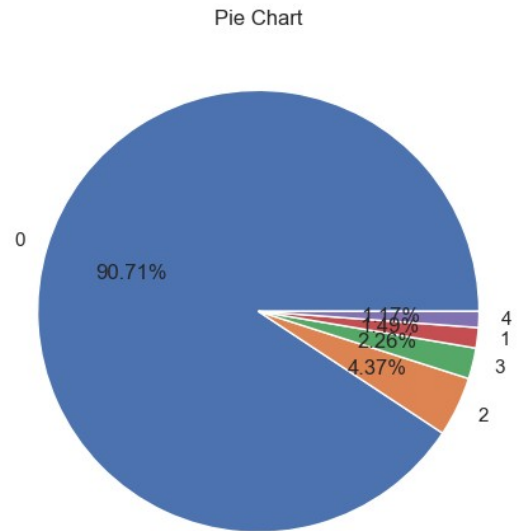

Floors



Condition



View






Define a Function for Barh Plot

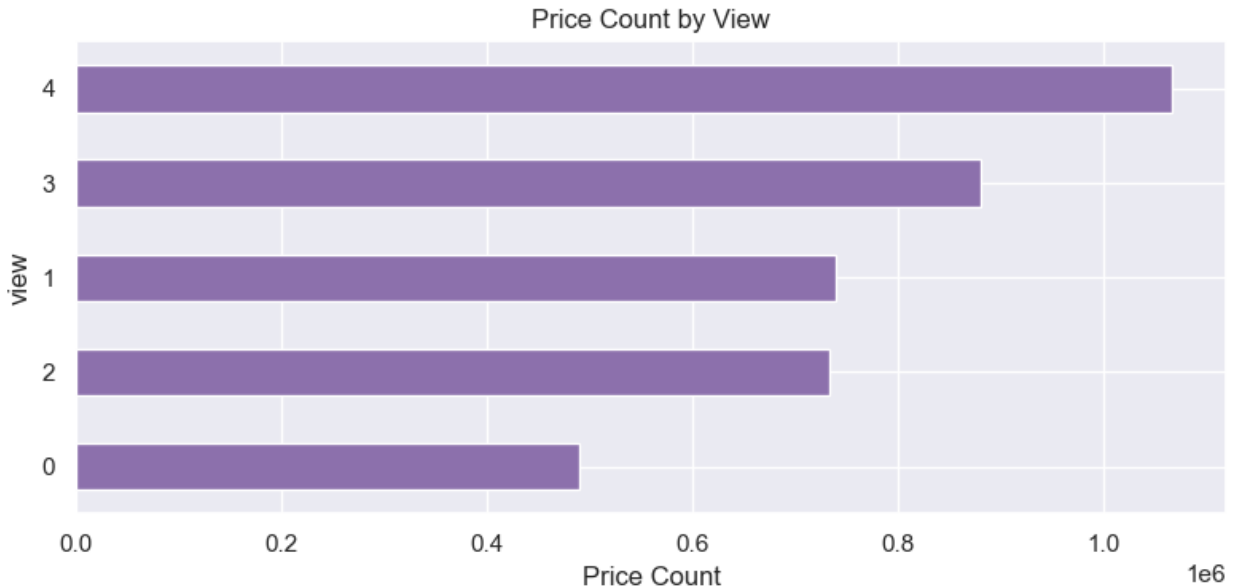
```
def bar_plot(x, y, df):
    barh = df.groupby([x])[y].mean()
    barh.sort_values(ascending=True, inplace=True)
    barh.plot(kind='barh', color = '#8c70ac', figsize=(8,4))
    plt.title(f'{title(y)} Count by {title(x)}')
    plt.xlabel(f'{title(y)} Count')
    plt.ylabel(x)

    plt.tight_layout()
    plt.show()
```

```
bar_plot('waterfront', 'price', df)
bar_plot('yr_renovated', 'price', df)
bar_plot('floors', 'price', df)
bar_plot('condition', 'price', df)
bar_plot('view', 'price', df)
```





Data Preprocessing

```
from sklearn.preprocessing import StandardScaler

# Initialize StandardScaler
stc = StandardScaler()

stc_cols = ['price', 'sqft_living', 'sqft_lot', 'sqft_above',
            'sqft_basement', 'yr_built', 'sqft_lot15', 'sqft_living15', 'long',
            'lat', ]

# Apply Standard Scaler to the selected columns
df[stc_cols] = stc.fit_transform(df[stc_cols])
```

Training and Evaluating Different Models

```
from sklearn.model_selection import train_test_split

x = df.drop(['price', 'id'], axis=1)
y = df['price']

x_train, x_test, y_train, y_test = train_test_split(x, y,
                                                    test_size=0.2, random_state=12)

#Importing the Libraries
from sklearn.ensemble import RandomForestRegressor,
GradientBoostingRegressor
from sklearn.linear_model import LinearRegression, Ridge
from sklearn.neighbors import KNeighborsRegressor
from sklearn.linear_model import LogisticRegression
from sklearn.datasets import make_classification
```

```

from sklearn.model_selection import GridSearchCV
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import VotingRegressor
from xgboost import XGBRegressor
from sklearn.metrics import mean_squared_error, r2_score

# List of Models to Try
models = [
    ('Linear Regression', LinearRegression()),
    ('Ridge Regression', Ridge()),
    ('Decision Tree', DecisionTreeRegressor()),
    ('Random Forest', RandomForestRegressor()),
    ('Gradient Boosting', GradientBoostingRegressor()),
    ('K-Nearest Neighbors', KNeighborsRegressor()),
]

# Train and evaluate each model
for name, model in models:
    model.fit(x_train, y_train)
    y_pred = model.predict(x_test)
    mse = mean_squared_error(y_test, y_pred)
    r2 = r2_score(y_test, y_pred)
    print(f'{name}: Mean Squared Error = {round(mse,3)}, R-squared = {round(r2, 3)}')

Linear Regression: Mean Squared Error = 0.279, R-squared = 0.723
Ridge Regression: Mean Squared Error = 0.279, R-squared = 0.723
Decision Tree: Mean Squared Error = 0.234, R-squared = 0.768
Random Forest: Mean Squared Error = 0.122, R-squared = 0.879
Gradient Boosting: Mean Squared Error = 0.132, R-squared = 0.869
K-Nearest Neighbors: Mean Squared Error = 0.194, R-squared = 0.808

rf_model_best = RandomForestRegressor(max_depth= 20, n_estimators=
200)
rf_model_best.fit(x_train, y_train)

# Predict using the updated features
y_pred_best = rf_model_best.predict(x_test)
mse_best = mean_squared_error(y_test, y_pred_best)
r2_best = r2_score(y_test, y_pred_best)


print(f'Mean Squared Error (Tuned Random Forest): {round(mse_best,
3)}')
print(f'R-squared (Tuned Random Forest): {round(r2_best, 3)}')

Mean Squared Error (Tuned Random Forest): 0.118
R-squared (Tuned Random Forest): 0.883

# Visualize the Predicted Prices Against the Actual Prices
plt.figure(figsize=(8, 6))
sns.scatterplot(x=y_test, y=y_pred_best)

```

```
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()],
linestyle='--', color='black', label='Perfect Prediction')
plt.title('Actual Prices vs. Predicted Prices (Tuned Random Forest)')
plt.ylabel('Predicted Prices')
plt.xlabel('Actual Prices')
plt.legend()
plt.show()
```



Summary and Conclusion for House Price Prediction Dataset

In this project, we aimed to predict house prices using a given dataset. The steps involved in the data preprocessing and model training are outlined below:

1. Date Column Transformation:
 - The date column was transformed into a more suitable format to ensure it could be effectively used in the model.
2. Reduction of Unique Values:
 - For certain columns, the number of unique values was reduced. This was done to simplify the model and improve its performance.

3. Removal of Outliers:
 - Outliers were identified and removed from the dataset to ensure the model was not adversely affected by extreme values.
4. Data Visualization:
 - Appropriate visualizations were created to understand the data distribution and relationships between different features. This step provided valuable insights into the dataset.
5. Standardization:
 - The data was standardized to ensure that all features contributed equally to the model training process. This involved scaling numerical features to a common range.
6. Model Training:
 - A Random Forest model was trained on the preprocessed data. This model was chosen for its robustness and ability to handle complex datasets.
7. Model Performance:
 - The trained Random Forest model achieved an accuracy of 88.3%. This indicates that the model performs well in predicting house prices based on the given features.

Conclusion

The steps taken in this project, from data preprocessing to model training, demonstrate a systematic approach to handling a house price prediction dataset. The transformation of the date column, reduction of unique values, removal of outliers, data visualization, and standardization were crucial in preparing the data for modeling. The Random Forest model proved to be effective, achieving a solid accuracy of 88.3%. This project highlights the importance of thorough data preprocessing and the application of a robust machine learning model in achieving accurate predictions.

This structured approach can be replicated and adapted for similar datasets to ensure effective data analysis and model performance.

Developed by Hosein Mohammadi

GitHub : <https://github.com/Hosein541>

Kaggle : <https://www.kaggle.com/hoseinnnnnnn>

Gmail : Huseinmohammadi83@gmail.com